

PROJECT

Object Classification

A part of the Deep Learning Nanodegree Foundation Program

PROJECT REVIEW

CODE REVIEW

NOTES

SHARE YOUR ACCOMPLISHMENT!  

Meets Specifications

Kudos ! I think you've done a perfect job of implementing a convolutional neural net fully. It's very clear that you have a good understanding of the basics. Keep improving and keep learning.

Advanced tips for improving net results

- Try and use deeper architectures, which have general tendency to blow up or vanish the gradients - so there's a net architecture known as Residual Nets, used to circumnavigate the issues with deeper architectures
- Try and use batch normalisation techniques. They basically normalise the output of every conv layer, very much for the same reason that you normalised the input image before feeding it to the net

If you are keen on learning a bit more into what Computer Vision Scientists use regularly in their nets. Try reading up a bit more on

- Batch Normalisation layers
 - Deconvolutional layers
 - Dilated Convolutional layers
- The details of all these layers are there in the TFLearn modules.

Keep up the good work !

Required Files and Tests



The project submission contains the project notebook, called "dlnd_image_classification.ipynb".



All the unit tests in project have passed.

Preprocessing



The `normalize` function normalizes image data in the range of 0 to 1, inclusive.

Good job in normalising the image intensities with global maximum value of 255 instead of using local image specific maxima



The `one_hot_encode` function encodes labels to one-hot encodings.

Good smart move in using `np.eye` directly for encoding !

Neural Network Layers

The neural net inputs functions have all returned the correct TF Placeholder.

Good job in implementing all the placeholders so perfectly !

The `conv2d_maxpool` function applies convolution and max pooling to a layer.

The convolutional layer should use a nonlinear activation.

This function shouldn't use any of the tensorflow functions in the `tf.contrib` or `tf.layers` namespace.

Very good job in implementing the conv, maxpool layers with the appropriate shape filters and adding the bias. This was one of the tougher challenges in the entire submission. Good you could solve it so easily :)

The `flatten` function flattens a tensor without affecting the batch size.

Appreciate that you used basic tensor operations for the reshape layer instead of using a direct off-the-shelf implementation, Very impressive indeed !

The `fully_conn` function creates a fully connected layer with a nonlinear activation.

Again, appreciate that you used basic tensor operations for the fully connected layer instead of using a direct off-the-shelf implementation, Very impressive, again !

The `output` function creates an output layer with a linear activation.

Again , very impressive and more so because of careful implementation of linear activation

Neural Network Architecture

The `conv_net` function creates a convolutional model and returns the logits. Dropout should be applied to alt least one layer.

looks like a perfect conv net to start off with !

Pro Tip : Given you've already implemented the net so well, there's some extra resources for you to read on. [This answer](#) is regarding how to choose the best architecture of convolutional layers.

Couple of rules of thumb to help you going ahead

- Try and always use batch_normalisation after conv layers before maxpool so that gradients don't overflow while training
- If you're downsampling image by factor of x (say using maxpool) , then always increase the number of filters in the subsequent convolution kernel by x i.e. if you're having `conv_num_outputs` as 16 and you maxpool by 2, then in the next conv layer increase `conv_num_outputs` to 32.

Neural Network Training

The `train_neural_network` function optimizes the neural network.

The `print_stats` function prints loss and validation accuracy.

The hyperparameters have been set to reasonable numbers.

Great job on tuning hyperparameters nicely !

Pro Tip : Given you've tuned your hyper parameters so well, you may want to read more about rational ways to tune all your hyper parameters by methods like random search, grid search etc. - http://neupy.com/2016/12/17/hyperparameter_optimization_for_neural_networks.html



The neural network validation and test accuracy are similar. Their accuracies are greater than 50%.

The accuracy for both test and validation accuracy is within acceptable ranges i.e. both are above 50% and within 10% of each other

 [DOWNLOAD PROJECT](#)

Have a question about your review? Email us at review-support@udacity.com and include the link to this review.

[RETURN TO PATH](#)

[Student FAQ](#)